

NAJOUNG KIM

najoung@bu.edu / najoung.kim

Office 138P, 111 Cummington Mall

Boston, MA, USA 02215

Updated August 2026

ACADEMIC POSITIONS	Boston University	2023–current
	Assistant Professor, Department of Linguistics	
	Affiliated Faculty, Department of Computer Science	
	Faculty Fellow, Center for Computing & Data Sciences	Until 2026
	New York University	2021–2022
	Faculty Fellow / Assistant Professor, Center for Data Science	
EDUCATION	Johns Hopkins University	2021
	Ph.D. Cognitive Science	
	Advisors: Paul Smolensky and Kyle Rawlins	
	Dissertation: Compositional Generalization in Artificial Neural Networks	
	Johns Hopkins University	2017
	M.A. Cognitive Science	
	University of Oxford	2015
	M.St. General Linguistics & Comparative Philology	Distinction
	Seoul National University	2014
	B.A. English Linguistics & Literature	Summa Cum Laude
	B.A. Linguistics (double major)	
	Minor: Computer Science & Engineering	
EXPERIENCE	Visiting Faculty Researcher, Google DeepMind	2024–2025
	Google Research	2022–2024
	<i>Host: Deepak Ramachandran → Chris Tar</i>	
	Reasoning and abstraction in large language models	
	Evaluation of large multimodal models	
	Student researcher, Google AI	Aug–Nov 2020
	Research intern, Google AI	May–Aug 2020
	<i>Hosts: Deepak Ramachandran, Ellie Pavlick</i>	
	Improving Question-Answering systems with presupposition verification	
	Research intern, IBM Research AI	May–Aug 2019
	<i>Hosts: Song Feng, Chulaka Gunasekara</i>	
	Evaluating sentence representations for discourse	
	Graduate student researcher,	Jun–Aug 2018
	Jelinek Memorial Summer Workshop (JSALT)	
	<i>PIs: Sam Bowman, Ellie Pavlick</i>	
	General-purpose Sentence Representation Learning team	

	Visiting researcher, Korea Advanced Institute of Science and Technology (KAIST) <i>PI: Jong C. Park</i> Computational modeling of semantic fluency	2015–2016
	Intern developer, NAVER Corporation Knowledge Extraction Team	2013–2014
AWARDS & FUNDING	Research partnership, MassMutual (\$80,000) <i>For Mechanistic understanding of reasoning and explanations in entity tracking</i>	2025–2026
	Cloudex Microgrant (~\$17,000 in H200 GPU support)	2026
	Best Paper Award, GenBench <i>For Is artificial intelligence still intelligence?</i> <i>LLMs generalize to novel adjective-noun pairs, but don't mimic the full human population</i>	2024
	National Deep Inference Fabric (NDIF) Pilot Program <i>For Mechanisms of entity tracking in language models (with Sebastian Schuster)</i>	2024
	Unrestricted gift, Google (\$60,000)	2024
	Benchmarking LLM agents on consequential real-world tasks Open Philanthropy (\$756,396) with Sebastian Schuster	2024–2026
	PaliGemma Academic Program , Google (\$5,000 in GCP credit) <i>Comparing linguistic representations learned by text-only vs. multimodal models</i> (with Kanishka Misra, Adam Dahlgren Lindström, Lucia Donatelli, & Yulu Qin)	2024
	CDS Faculty Fellowship Center for Data Science, Boston University (\$50,000)	2023–2026
	Area Chair Award (Interpretability and Analysis Track), ACL 2023 <i>For Entity Tracking in Language Models</i>	2023
	NSF Doctoral Dissertation Research Improvement Grant: <i>Compositional Linguistic Generalization in Human and Machine Learning</i> BCS-2041221 (\$12,771) with Paul Smolensky, Geraldine Legendre & Tal Linzen.	2021–2022
	NeurIPS 2019 Travel Award	2019
	Best Paper Award, *SEM 2019 <i>For Probing what different NLP tasks teach machines about function word comprehension</i>	2019
	Owen Scholarship Awarded to select incoming doctoral students in natural sciences (\$18,000)	2016–2019
	Mica and Ahmet Ertegun Graduate Scholarship Programme in the Humanities Full scholarship for Master's degree at University of Oxford (~\$42,000)	2014–2015
	SNU Undergraduate Research Program Grant for select undergraduate research proposals (~\$3000)	2013

SNU Foundation Funds Scholarship
SNU Superior Academic Performance Scholarship

2013
2010–2011

WORK IN PROGRESS

Jane Li and **Najoung Kim**, [Linear representations of grammaticality in large language models](#). arXiv.

Najoung Kim, [No Escape from Behavior in Evaluating Compositionality](#). PsyArXiv.

Xiang Fu,* Seungmin Cho,* Yukyung Lee, and **Najoung Kim**. LLMs Learn Better In-Context from Rules than from Examples. In submission. (*Equal contribution)

Kanishka Misra and **Najoung Kim**. [Generating novel experimental hypotheses from language models: A case study on cross-dative generalization](#). In submission.

PUBLICATIONS

Jing Liu and **Najoung Kim**. Does Episodic Memory Help Close the Lexical Frequency Gap in Sensitivity to Syntactic Contrasts? A Test Using Retrieval-Augmented Language Models (2026). Conference on Language Modeling (COLM).

Zilu Tang,* Qiao Zhao,* Gabriel Franco, Derry Tanti Wijaya, Aaron Mueller, Sebastian Schuster, and **Najoung Kim** (2026). [Do Language Models Track Entities Across State Changes?](#) *Forty-third International Conference on Machine Learning (ICML)*. (*Equal contribution)

Nicholas Edwards,* Yukyung Lee,* Yujun Audrey Mao, Yulu Qin, Sebastian Schuster,[†] **Najoung Kim**[†] (2026). [RExBench: Can coding agents autonomously implement AI research extensions?](#) *Annual Conference of the Association for Computational Linguistics (ACL)*. (*,[†]Equal contribution)

Eunsu Kim, Junyeong Park, Juhyun Oh, Kiwoong Park, Seyoung Song, A. Seza Doğruöz, Alice Oh, and **Najoung Kim** (2026). [Are they lovers or friends? Evaluating LLMs’ Social Reasoning in English and Korean Dialogues](#). *Annual Conference of the Association for Computational Linguistics (ACL)*.

Arkadiy Saakyan, **Najoung Kim**, Smaranda Muresan, and Tuhin Chakrabarty (2026). [Death of the Novel\(ty\): Beyond N-Gram Novelty as a Metric for Textual Creativity](#) *International Conference on Learning Representations (ICLR)*.

Yulu Qin,* Dheeraj Varghese,* Adam Dahlgren Lindström, Lucia Donatelli, Kanishka Misra,[†] and **Najoung Kim**[†] (2025). [Vision-and-Language Training Helps Deploy Taxonomic Knowledge but Does Not Fundamentally Alter It](#). *The Thirty-Ninth Annual Conference on Neural Information Processing Systems (NeurIPS)* (*,[†]Equal contribution)

Yukyung Lee, JoongHoon Kim, Jaehee Kim, Hyowon Cho, Jaewook Kang, Pilsung Kang, and **Najoung Kim** (2025). [CheckEval: A reliable LLM-as-a-Judge framework for evaluating text generation using checklists](#). *Empirical Methods in Natural Language Processing (EMNLP)*. [Nominated for Outstanding Paper](#)

Hayley Ross, Kathryn Davidson, and **Najoung Kim** (2025). [Is analogy enough to draw novel adjective-noun inferences?](#) *Society for Computation in Linguistics (SCiL)*. [Selected as talk](#)

Abulhair Saparov, Srushti Pawar, Shreyas Pimpalgaonkar, Nitish Joshi, Richard Yuanzhe

Pang, Vishakh Padmakumar, Seyed Mehran Kazemi, **Najoung Kim**,* and He He* (2025). [Transformers Struggle to Learn to Search Without In-context Exploration](#). *International Conference on Learning Representations (ICLR)*. (*Equal contribution)

Hayley Ross, **Najoung Kim**, and Kathryn Davidson (2025). [Fake reefs are sometimes reefs and sometimes not, but are always compositional](#). *Experiments in Linguistic Meaning (ELM)* 3. [Selected as talk](#)

Aditya Yedetore and **Najoung Kim** (2024). [Semantic Training Signals Promote Hierarchical Syntactic Generalization in Transformers](#). *Empirical Methods in Natural Language Processing (EMNLP)*.

Nitish Joshi, Javier Rando, Abulhair Saparov, **Najoung Kim**, and He He (2024). [Personas as a Way to Model Truthfulness in Language Models](#). *Empirical Methods in Natural Language Processing (EMNLP)*.

Hayley Ross, Kathryn Davidson, and **Najoung Kim** (2024). [Is artificial intelligence still intelligence? LLMs generalize to novel adjective-noun pairs, but don't mimic the full human population](#). *GenBench Workshop @ EMNLP*. **Best Paper Award**

Najoung Kim,* Sebastian Schuster,* and Shubham Toshniwal* (2024). [Code Pre-training Improves Entity Tracking Abilities of Language Models](#). arXiv. (*Equal contribution)

Najoung Kim and Paul Smolensky (2024). [Structural Generalization of Modification in Adult Learners of an Artificial Language](#). *Proceedings of the Annual Meeting of the Cognitive Science Society (CogSci)*. [Selected as talk](#)

Katherine M. Collins, **Najoung Kim**, Yonatan Bitton, Verena Rieser, Shayegan Omidshafiei, Yushi Hu, Sherol Chen, Senjuti Dutta, Minsuk Chang, Kimin Lee, Youwei Liang, Georgina Evans, Sahil Singla, Gang Li, Adrian Weller, Junfeng He, Deepak Ramachandran, and Krishnamurthy Dj Dvijotham (2024). [Beyond Thumbs Up/Down: Untangling Challenges of Fine-Grained Feedback for Text-to-Image Generation](#). *The Seventh AAAI/ACM Conference on AI, Ethics, and Society (AIES)*.

Ashwin Daswani, Rohan Sawant, and **Najoung Kim** (2024). [Syn-\(QA\)²: Evaluating False Assumptions in Long-tail Questions with Synthetic QA Datasets](#). arXiv.

Zhaofeng Wu, Linlu Qiu, Alexis Ross, Ekin Akyürek, Boyuan Chen, Bailin Wang, **Najoung Kim**, Jacob Andreas, and Yoon Kim (2024). [Reasoning or Reciting? Exploring the Capabilities and Limitations of Language Models Through Counterfactual Tasks](#). *North American Chapter of the Association for Computational Linguistics (NAACL)*.

Bingzhi Li, Lucia Donatelli, Alexander Koller, Tal Linzen, Yuekun Yao, and **Najoung Kim** (2023). [SLOG: A Structural Generalization Benchmark for Semantic Parsing](#). *Empirical Methods in Natural Language Processing (EMNLP)*.

Jason Wei,* **Najoung Kim**,* Yi Tay, and Quoc V. Le (2023). [Inverse scaling can become U-shaped](#). *Empirical Methods in Natural Language Processing (EMNLP)*. (*Equal contribution)

Najoung Kim* and Sebastian Schuster* (2023). [Entity Tracking in Language Models](#). *Annual Conference of the Association for Computational Linguistics (ACL)*. (*Equal contribution)

contribution) **Area Chair Award**

Najoung Kim*, Phu Mon Htut*, Samuel R. Bowman, and Jackson Petty (2023). [\(QA\)²: Question Answering with Questionable Assumptions](#). *Annual Conference of the Association for Computational Linguistics (ACL)*. (*Equal contribution)

Seyed Mehran Kazemi, **Najoung Kim**, Deepti Bhatia, Xin Xu, and Deepak Ramachandran (2023). [LAMBADA: Backward Chaining for Automated Reasoning in Natural Language](#). *Annual Conference of the Association for Computational Linguistics (ACL)*.

Najoung Kim, Jatin Khilnani, Alex Warstadt, and Abed Qaddoumi (2023). [Reconstruction Probing](#). *Findings of the Annual Conference of the Association for Computational Linguistics (Findings of ACL)*.

Abulhair Saparov, Richard Yuanzhe Pang, Vishakh Padmakumar, Nitish Joshi, Seyed Mehran Kazemi, **Najoung Kim**,* and He He* (2023). [Testing the General Deductive Reasoning Capacity of Large Language Models Using OOD Examples](#). *Conference on Neural Information Processing Systems (NeurIPS)*. (*Equal contribution)

Mehran Kazemi, Quan Yuan, Deepti Bhatia, **Najoung Kim**, Xin Xu, Vaiva Imbraisaite, and Deepak Ramachandran (2023). [BoardgameQA: A Dataset for Natural Language Reasoning with Contradictory Information](#). *Conference on Neural Information Processing Systems (NeurIPS), Datasets and Benchmark Track*.

Ian McKenzie, Alexander Lyzhov, Michael Pieler, Alicia Parrish, Ameya Prabhu, Aaron Mueller, Euan McLean, [Aaron Kirtland](#), [Alexis Ross](#), [Alisa Liu](#), [Andrew Gritsevskiy](#), [Daniel Wurgaft](#), [Derik Kauffman](#), [Gabriel Recchia](#), [Jiacheng Liu](#), [Joe Cavanagh](#), [Max Weiss](#), [Sicong Huang](#), [The Floating Droid](#), [Tom Tseng](#), [Tomasz Korbak](#), [Xudong Shen](#), [Yuhui Zhang](#), [Zhengping Zhou*](#), **Najoung Kim**, Samuel R. Bowman, and Ethan Perez (2023). Inverse Scaling: When Bigger Isn't Better. *Transactions on Machine Learning Research (TMLR)*. (*Winning task authors) **Featured Certification**

Wentao Wang, Wai Keen Vong, **Najoung Kim**, and Brenden Lake (2023). [Finding Structure in One Child's Linguistic Experience](#). *Cognitive Science*.

Najoung Kim, Tal Linzen, and Paul Smolensky (2022). [Uncontrolled Lexical Exposure Leads to Overestimation of Compositional Generalization in Pretrained Models](#). arXiv.

Najoung Kim and Paul Smolensky (2021). [Testing for Grammatical Category Abstraction in Neural Language Models](#). *The Society for Computation in Linguistics (SCiL)*. Selected as talk

Najoung Kim, Ellie Pavlick, Burcu Karagol Ayan, and Deepak Ramachandran (2021). [Which Linguist Invented the Lightbulb? Presupposition Verification for Question-Answering](#). *Annual Conference of the Association for Computational Linguistics (ACL)*.

Najoung Kim and Tal Linzen (2020). [COGS: A Compositional Generalization Challenge Based on Semantic Interpretation](#). *Empirical Methods in Natural Language Processing (EMNLP)*.

Najoung Kim, Song Feng, Chulaka Gunasekara, and Luis A. Lastras (2020). [Implicit Discourse Relation Classification: We Need to Talk About Evaluation](#). *Annual Conference of the Association for Computational Linguistics (ACL)*.

Najoung Kim and Tal Linzen (2019). [Compositionality as directional consistency in sequential neural networks](#). *Context and Compositionality in Biological and Artificial Neural Systems Workshop, NeurIPS*.

Najoung Kim, Roma Patel, Adam Poliak, Alex Wang, Patrick Xia, Tom McCoy, Ian Tenney, Alexis Ross, Tal Linzen, and Benjamin Van Durme, Samuel R. Bowman, and Ellie Pavlick (2019). [Probing what different NLP tasks teach machines about function word comprehension](#). *Proceedings of the Eighth Joint Conference on Lexical and Computational Semantics (*SEM)*. **Best Paper Award**

Alex Wang, Jan Hula, Patrick Xia, Raghavendra Pappagari, R. Thomas McCoy, Roma Patel, **Najoung Kim**, Ian Tenney, Yinghui Huang, Katherin Yu, Shuning Jin, Berlin Chen, Benjamin Van Durme, Edouard Grave, Ellie Pavlick, and Samuel R. Bowman (2019). [How to Get Past Sesame Street: Sentence-Level Pretraining Beyond Language Modeling](#). *Proceedings of the Annual Conference of the Association for Computational Linguistics (ACL)*.

Najoung Kim, Jung-Ho Kim, Maria K. Wolters, Sarah E. MacPherson, and Jong C. Park (2019). [Automatic Scoring of Semantic Fluency](#). *Frontiers in Psychology*.

Najoung Kim, Kyle Rawlins, Benjamin Van Durme, and Paul Smolensky (2019). [Predicting the Argumenthood of English Prepositional Phrases](#). *Proceedings of the 33rd AAAI Conference on Artificial Intelligence (AAAI)*.

Ian Tenney, Patrick Xia, Berlin Chen, Alex Wang, Adam Poliak, R. Thomas McCoy, **Najoung Kim**, Benjamin Van Durme, Samuel R. Bowman, Dipanjan Das, and Ellie Pavlick (2019). [What do you learn from context? Probing for sentence structure in contextualized word representations](#). *International Conference on Learning Representations (ICLR)*.

Maria K. Wolters, **Najoung Kim**, Jung-Ho Kim, Sarah E. MacPherson, and Jong C. Park (2016). [Prosodic and Linguistic Analysis of Semantic Fluency Data: A Window into Speech Production and Cognition](#). *Interspeech*.

Jung-Ho Kim, **Najoung Kim**, Hanchool Park, and Jong C. Park (2016). Enhanced Sign Language Transcription System via Hand Tracking and Pose Estimation. *Journal of Computing Science and Engineering* vol 10.3.

Najoung Kim and Jong C. Park (2016). [A Morphological Approach to the Longitudinal Detection of Dementia](#). *Proceedings of HCI Korea 2016, The HCI Society of Korea*.

PRESENTATIONS WITHOUT PROCEEDINGS

Jing Liu and **Najoung Kim** (2026). Does Episodic Memory Help Close the Lexical Frequency Gap in Sensitivity to Syntactic Contrasts? A Test Using Retrieval-Augmented Language Models. *The Annual Meeting of the Cognitive Science Society (CogSci)*.

Zilu (Peter) Tang, Qiao Zhao, Gabriel Franco, Geneva Yang, Angelos Poulis, Derry Wijaya, Aaron Mueller, Sebastian Schuster, and **Najoung Kim** (2025). Mechanistic Understanding of Entity Tracking in Natural Language involving Multiple Operations. *New England Mechanistic Interpretability Workshop (NEMI)*.

Aditya Yedetore and **Najoung Kim** (2025). Implicit mechanisms for symbol manipulation in RNNs. *New England NLP Meeting Series (NENLP)*.

Hayley Ross, Kathryn Davidson, and **Najoung Kim** (2025). Is analogy enough to draw novel adjective-noun inferences? *New England NLP Meeting Series (NENLP)*.

Aditya Yedetore and **Najoung Kim** (2024). Semantic Training Signals Promote Hierarchical Syntactic Generalization in Neural Networks.

- *The 49th Boston University Conference on Language Development (BUCLD)*.
- *New England NLP Meeting Series (NENLP)*. [Selected as spotlight talk](#)

Hayley Ross, Kathryn Davidson, and **Najoung Kim** (2024). When is a fake concert still a concert? A study of adjective-noun composition in LLMs. *New England NLP Meeting Series (NENLP)*.

Kanishka Misra and **Najoung Kim** (2023). [Abstraction via exemplars? A representational case study on lexical category inference in BERT](#). Presented at *The 48th Boston University Conference on Language Development (BUCLD)*.

Sadhwi Srinivas,* **Najoung Kim**,* and Kyle Rawlins (2020). [Maximize Presupposition! and the Korean demonstrative ku](#). Presented at *The 94th Annual Meeting of the Linguistic Society of America (LSA)*. [[poster](#)] (*Equal contribution)

Najoung Kim, Benjamin Van Durme, Ellie Pavlick, and Paul Smolensky (2018). Linguistically informed tasks for evaluating structure encoded by sentence representations. *WeCNLP Summit*, Facebook HQ, Menlo Park, CA. [Selected as spotlight talk](#)

Najoung Kim, Kyle Rawlins, and Paul Smolensky (2018). A gradient blend analysis of English PP verbal dependents. *Conference on Interdisciplinary Approaches to Linguistic Theory (CiALT) 2*, Humboldt-Universität zu Berlin.

Najoung Kim, Kyle Rawlins, and Paul Smolensky (2018). A gradient blend analysis of English PP verbal dependents. *Acceptability judgments in current linguistic theory*, Universitat Autònoma de Barcelona.

OTHER ARTICLES

Yonatan Belinkov, Aaron Mueller, **Najoung Kim**, Hosein Mohebbi Hanjie Chen, Dana Arad, and Gabriele Sarti (2025). [Proceedings of the 8th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP](#). Association for Computational Linguistics.

Arkadiy Saakyan, Josh Lee, Michal Todorovic, Deepak Ramachandran, Quan Yuan, Isabelle Guyon, and **Najoung Kim** (2025). Evaluating Critic Models for Human-AI Co-Creation: A Case Study with AI Critiques of Presentation Slides. Unpublished manuscript available upon request.

Yonatan Belinkov, **Najoung Kim**, Jaap Jumelet, Hosein Mohebbi, Aaron Mueller, and Hanjie Chen (2024). [Proceedings of the 7th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP](#). Association for Computational Linguistics.

Yonatan Belinkov, Sophie Hao, Jaap Jumelet, **Najoung Kim**, Arya McCarthy, and Hosein Mohebbi (2023). [Proceedings of the 6th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP](#). Association for Computational Linguistics.

Teven Le Scao, Angela Fan, ... **Najoung Kim** (Evaluation and Interpretability) ...

(2022), [BLOOM: A 176B-Parameter Open-Access Multilingual Language Model](#). arXiv.

Najoung Kim (2022). [Compositional Generalization in Artificial Neural Networks](#). PhD Dissertation, Johns Hopkins University. Committee: Paul Smolensky, Kyle Rawlins, Benjamin Van Durme, Tal Linzen, and Bob Frank.

SELECTED TALKS / TUTORIALS Invited talk at [Yale Linguistics Colloquium Series](#), Forthcoming in Fall 2026.

How do we move beyond associative insights from Artificial Intelligence?

- Invited talk at Linguistic Theory at the Age of LLMs - CreteLing Summer School. July 2026.
- Invited talk at Brown CoPsy. March 2026.

Language Models for Language Sciences. Invited talk at [Stanford NLP Seminar](#). May 2026.

(How) Do Language Models Track Entities Across State Changes? Invited talk at R&D Knowledge Sharing, MassMutual. March 2026.

Classical Computation in Connectionist Models. Invited talk at the [Cognitive Basis of Reasoning \(in Minds and AI\) Workshop](#), IVADO. Jan 2026.

Whence insights? Delineating human and machine CogSci, and hypothesis generation as a potential bridge. Invited talk at UCSD Linguistics. Dec 2025.

Whence insights? The value of delineating human and machine CogSci. Plenary address at the [Society for Language Development \(SLD\) Annual Symposium](#). Nov 2025.

RExBench: Can coding agents autonomously implement AI research extensions? Invited talk at IBM Research, HCAI Speaker Series. August 2025.

What does it take to convince ourselves that a system is exhibiting compositionality?

- Keynote at [Copenhagen NLP Symposium](#). June 2025.
- Keynote at the [10th Workshop on Representation Learning for NLP \(RepL4NLP\)](#), co-located with NAACL 2025. May 2025.

Evaluating Multimodal Critic Models for Human-AI Co-Creation. Keynote at the [Structure and Generalization in Multimodal Language Understanding \(SAGE-MLU\) Workshop](#). March 2025.

Semantic generalizations in humans and machines.

- Keynote at the [Workshop on Generalization in NLP \(GenBench\) 2024](#) (co-located with EMNLP 2024). November 2024.
- Invited talk at Sungkyunkwan University, December 2024.

Generating novel experimental hypotheses from language models: A case study on cross-dative generalization. Invited talk at the [Language & Cognition Seminar at Harvard \(LangCog\)](#). October 2024.

Human and Machine Inductive Biases for Compositional Linguistic Generalization. Invited talk at the [Emerging Generalization Settings Workshop at Simons Institute](#). September 2024.

Comparing human and machine inductive biases for compositional linguistic generalization using semantic parsing. Invited talk at the [Fifth International Convention on the Mathematics of Neuroscience and AI \(Neuromonster\)](#). May 2024.

Linguistic tests as unit tests for AI systems. Invited talk at [New Horizons in Language Science: Large Language Models, Language Structure, and the Cognitive and Neural Basis of Language](#). National Science Foundation (NSF). May 2024.

Entity Tracking in Language Models.

- Invited talk at [UMass NLP](#). November 2024.
- Invited talk at KAIST Electrical Engineering. July 2024.
- Invited talk at AIR seminar at Boston University. April 2024.
- Invited talk at [South by Semantics Workshop at UT Austin](#). Feb 2024.
- Invited talk at [Penn NLP CLunch Seminar](#). Oct 2023.

Can we use simulated learners to prune the hypothesis space for human experiments?

- Invited talk at CPL Lab, MIT. April 2024.
- Invited talk at [Brown Linguistics](#). March 2024.

Inverse scaling: mitigation strategies and open questions. Keynote talk at the [Workshop on the Scaling Behavior of Large Language Models \(SCALE-LLM\)](#). March 2024.

Challenges in logical reasoning with LLMs. Invited talk at [Yale/Google ML Workshop: Theory and Practice of Foundation Models](#). Oct 2023.

Evidence for abstraction-via-exemplars from lexical category inference in neural language models. Invited talk at the [Virtual Psycholinguistics Forum](#). Sep 2023.

Compositional Linguistic Generalization in Artificial Neural Networks: Taking Stock. Invited talk at [Harvard Universals Workshop](#), Apr 2023.

Questions with Questionable Assumptions as a Challenge to QA Systems. Invited talk at Colgate University. Mar 2023.

Compositional Linguistic Generalization in Artificial Neural Networks: Taking Stock.

- Invited talk at Seminars on Formal Languages and Neural Networks (FlaNN), Nov 2022.
- Invited talk at CUNY Graduate Center Computational Linguistics Talk Series, Nov 2022.

Compositional Linguistic Generalization in Artificial Neural Networks.

- Invited talk at CompLing Lab, University of Chicago, May 2022.
- ‘Tech Talk’ at Google, May 2022.
- Invited talk at Boston University Linguistics Colloquium Series, Mar 2022.
- Invited talk at [MIT Complang](#). Mar 2022.
- Invited talk at [Seminar on the interactions between formal and computational linguistics \(ILFC\)](#). Feb 2022.
- Invited talk at Linguistics Colloquium, Seoul National University. Dec 2021.

- Invited talk at [KAIST Electrical Engineering Colloquium Lecture Series](#). Nov 2021.
- Invited talk at [NYU ConCats](#). Nov 2021.
- Invited talk at the University of Michigan Cognitive Science Community Colloquium. Oct 2021.
- Invited talk at the [New York Philosophy of Language Workshop](#). Oct 2021.
- Invited talk at NYU Center for Data Science. Sep 2021.

Compositional Linguistic Generalization in Contemporary Neural Models of Language. Invited talk at Cornell Computational Psycholinguistics Group. Apr 2021.

What Aspects of Meaning are Missing from Current Natural Language Understanding Systems? Invited talk at Boston University, Mar 2021.

COGS: A Compositional Generalization Challenge Based on Semantic Interpretation.

- Invited talk at NERT lab, Georgetown University, Nov 2020.
- Invited poster presentation at the [Microsoft Research AI Breakthroughs Workshop](#), Sep 2020.

Probing what different NLP tasks teach machines about function word comprehension and where to go next. Invited talk at JHU CLSP Seminar, Oct 2019.

The complement-adjunct distinction as gradient blends: the case of English prepositional phrases. Invited talk at the Gradient Symbolic Computation Workshop, Baltimore, Sep 2019.

Semantic Role Labeling Tutorial (with Diego Marcheggiani, Michael Roth, and Benjamin Van Durme), EMNLP, Copenhagen, Sep 2017.

Approximating the Semantic Structures behind Category Fluency Sequences (Poster). MACSIM 6 at CUNY, New York, Oct 2016.

Detection and Categorisation of Neograms in Korean Text (+a). Invited talk at Ertegun House, Oxford, 2015.

MEDIA COVERAGE

[Linguistics in the Age of Artificial Intelligence](#). Boston University Arts & Sciences Magazine, April 2024.

[Compositional Generalization in Neural Networks](#). Interview at NLP Highlights Podcast hosted by Allen Institute for AI, Jan 2023.

ADVISING / MENTORING

Postdoctoral Fellows:

- [Yukyung Lee](#), 2024–

PhD Advisees:

- [Yulu Qin](#) (BU Computing & Data Sciences), 2024–
- [Audrey Mao](#) (BU Computer Science), 2024–
- [Hayley Ross](#) (Harvard Linguistics, co-advised with Kathryn Davidson), 2023–2025 → [NVIDIA](#)

- [Jing Liu](#) (École Normale Supérieure, co-advised with Emmanuel Dupoux & David Harwarth), 2023–
- [Aditya Yedetore](#) (BU Linguistics), 2022–

PhD Committees:

- Apoorv Khandelwal (Dissertation, Brown CS), 2026
- Xiulin Yang (Dissertation, Georgetown Linguistics), 2025–
- Ousmane Cisse (Dissertation prospectus, BU Linguistics), 2025–
- Cheng Tang (Dissertation, MIT Brain & Cognitive Science), 2025–
- Afra Feyza Akyürek (Dissertation & Prospectus, BU Computer Science), 2024–2025
- Yusuf Kocyigit (Dissertation & Qualifying exam, BU Computer Science), 2025
- Gabriel Franco (Qualifying exam, BU Computer Science), 2025
- Zhongping Zhang (Dissertation & Prospectus, BU Computer Science), 2024
- Andrea Burns (Dissertation, BU Computer Science), 2023
- Isidora Tourni (Dissertation, BU Computer Science), 2023

Other Ad-hoc Advising & Mentoring:

- Angelo Amato, Undergraduate researcher @ BU, 2025–current.
 - Supported by [BU Undergraduate Research Opportunities Program \(UROP\)](#), Summer 26
- Xiang Fu, Undergraduate researcher @ BU, 2025–current.
 - Supported by [BU Undergraduate Research Opportunities Program \(UROP\)](#), F25 & S26
- Nicholas Edwards (co-mentored with Sebastian Schuster), Research Assistant at University College London. 2024–current.
 - Outcome: Paper at ACL 2026
- Arkadiy Saakyan (co-hosted with Deepak Ramachandran), Student Researcher at Google DeepMind, May–Aug 2024.
 - Outcome: Internal report (unpublished due to company policy), return internship offer, subsequent collaboration leading to a paper at ICLR 2026
- Alyssa Feinberg (undergraduate at Brown University), Research Assistant @ BU, Summer 2024.
 - Outcome: Contribution to Code Pretraining manuscript
- Alara Balcisoy, Undergraduate researcher @ BU, Spring 2024.
 - Supported by [BU Undergraduate Research Opportunities Program \(UROP\)](#)
 - Outcome: Recipient of [BUCH Award](#)
- Bingzhi Li (co-advised with Tal Linzen), Visiting PhD student from Université de Paris, Fall 2022.
 - Outcome: Paper at EMNLP
- Abed Qaddoumi (co-mentored with Alex Warstadt), MS at New York University. 2021–2022. *Current: PhD student at Stony Brook Linguistics*

- Outcome: Paper at ACL Findings
- Jatin Khilnani (co-mentored with Alex Warstadt), MS at New York University. 2021–2022. *Current: PhD student at University of Pittsburgh CS*
 - Outcome: Paper at ACL Findings
- Pablo Santos (co-mentored with Phu Mon Htut), [CDS-Courant Undergraduate Research Program \(CURP\)](#), Spring 2022.

TEACHING

Boston University:

- Fall 2025, Topics in Linguistics: Contemporary Research in Neural Network Models of Language
- Fall 2024–2025, Topics in Linguistics: Cognitive Science of Language
- Spring 2024–2025, Computational Linguistics
- Fall 2023–2024, Introduction to Programming for Computational Linguistics
- Spring 2023–2024, Topics in Linguistics: Metrics and Evaluation in Natural Language Processing

New York University:

- Fall 2021–2022, Capstone Project in Data Science

Johns Hopkins University:

- Fall 2019, Introduction to Computational Cognitive Science (TA/Co-Instructor)
- Spring 2019, Foundations of Neural Networks (TA/Lab Instructor)
- Spring 2018, Foundations of Cognitive Science (TA)
- Fall 2017, Semantics I (TA/Lab Instructor)
- Spring 2017, Language and Advertising (TA)

Seoul National University:

- Winter 2014, Samsung Convergence Software Course (Instructor)

PROFESSIONAL SERVICE

Grant panels:

- Advisory Panel, US National Science Foundation (NSF) (2024)

Editorial Team:

- Standing reviewer, Transactions of the Association for Computational Linguistics (2025–)

Conference Chairs:

- COLM 2026 (Area Chair)
- ACL 2026 (Website Chair)
- EMNLP 2025 (Area Chair, Best Paper Committee)
- COLM 2025 (Area Chair)
- ACL 2025 (Senior Area Chair for Linguistic Theories, Cognitive Modeling, and Psycholinguistics)
- EMNLP 2024 (Senior Area Chair for Linguistic Theories, Cognitive Modeling, and Psycholinguistics)

- NAACL 2024 (Publicity Chair)
- ACL 2023 (Senior Area Chair for Linguistic Theories, Cognitive Modeling, and Psycholinguistics)
- *SEM 2022 (Area Chair for Psycholinguistics, cognitive linguistics and semantic processing)

Organization:

- [New England Mechanistic Interpretability Workshop](#) (2025)
- [BlackboxNLP 2025](#) (Co-located with EMNLP 2025)
- [BlackboxNLP 2024](#) (Co-located with EMNLP 2024)
- AI4Research (Co-located with IJCAI 2024)
- [BlackboxNLP 2023](#) (Co-located with EMNLP 2023)
- [Inverse Scaling Prize](#) (2022)

Ad-hoc journal reviews:

- Transactions of the Association for Computational Linguistics (2025)
- Journal of Memory and Language (2024)
- Computational Linguistics (2024)
- Language, Cognition and Neuroscience (2024)
- Cognitive Science (2023, 2022)
- Nature Machine Intelligence (2023, 2022)

Ad-hoc grant reviews:

- Israel Science Foundation (2024)

Conference/Workshop reviews:

- ICLR (2026*)
- NAACL Student Research Workshop (participated as mentor, 2025)
- COLM (2024)
- *SEM (2024, 2023, * 2022, 2019)
- Scale-LLM Workshop (2024)
- EACL Student Research Workshop (2023, 2021)
- EMNLP (2022, 2020, * 2018)
- COLING (2022)
- Bridges and Gaps between Formal and Computational Linguistics (Workshop at ESSLI 2022)
- CogSci (2022)
- ACL Student Research Workshop, (2019–2022)
- ACL-IJCNLP (2021)
- ACL (2020)
- EMNLP-IJCNLP (2019)

* Acknowledged as outstanding reviewer (or equivalent)

PROFESSIONAL MEMBERSHIP	Association for Computational Linguistics Cognitive Science Society Linguistic Society of America	
COMPUTER SKILLS	Python, R (proficient) Java, C, Scheme/Racket, OCaml (basic) Tools: PyTorch, PCIBex, Mechanical Turk, Prolific	
HUMAN LANGUAGES	Korean (native); English (near-native); Spanish (intermediate); Japanese (reading knowledge)	
COMMUNITY ENGAGEMENT	Panel , Brandeis CL/NLP Industry Event	2026
	Panel , New England Mechanistic Interpretability Workshop (NEMI)	2025
	Dagstuhl Seminar , Linguistics and Language Models: What Can They Learn from Each Other?	2025
	Panel , Workshop on Representation Learning for NLP (RepL4NLP)	2025
	Panel , New England NLP Meeting Series (NENLP)	2025
	Panel , Learning to Think after ChatGPT: A Panel Discussion <i>Center for Data Science, Boston University</i>	2023
	Panel , KASELL Fall Conference	2021
	Diversity & Representation Committee	2021
	Colloquium Student Committee <i>Department of Cognitive Science, Johns Hopkins University</i>	2018–2021
	Panel , BrainIAC Professional Development Event	9/21/2020
	Organizer , Corpus Tools: Past, Present, and Future. <i>Ertegun House, Oxford</i>	2015
	Mentor, Samsung Convergence Software Course <i>Department of Computer Science, Seoul National University</i>	2014
	Peer tutor, English (volunteer work) <i>Seoul National University</i>	2011
	English tutor (volunteer work) <i>Younghak High School</i>	2010