

NAJOUNG KIM

najoung@bu.edu / najoung.kim

Office 808, 665 Commonwealth Ave.

Boston, MA, USA 02215

Updated August 2025

ACADEMIC POSITIONS	Boston University	2023–current
	Assistant Professor, Department of Linguistics	
	Affiliated Faculty, Department of Computer Science	
	Faculty Fellow, Center for Computing & Data Sciences	Until 2026
	New York University	2021–2022
	Faculty Fellow / Assistant Professor, Center for Data Science	
EDUCATION	Johns Hopkins University	2021
	Ph.D. Cognitive Science	
	Advisors: Paul Smolensky and Kyle Rawlins	
	Johns Hopkins University	2017
	M.A. Cognitive Science	
	University of Oxford	2015
	M.St. General Linguistics & Comparative Philology	Distinction
	Seoul National University	2014
	B.A. English Linguistics & Literature	Summa Cum Laude
	B.A. Linguistics (double major)	
	Minor: Computer Science & Engineering	
EXPERIENCE	Visiting Faculty Researcher, Google DeepMind	2024–2025
	Google Research	2022–2024
	<i>Host: Deepak Ramachandran → Chris Tar</i>	
	Reasoning and abstraction in large language models	
	Evaluation of large multimodal models	
	Student researcher, Google AI	Aug–Nov 2020
	Research intern, Google AI	May–Aug 2020
	<i>Hosts: Deepak Ramachandran, Ellie Pavlick</i>	
	Improving Question-Answering systems with presupposition verification	
	Research intern, IBM Research AI	May–Aug 2019
	<i>Hosts: Song Feng, Chulaka Gunasekara</i>	
	Evaluating sentence representations for discourse	
	Graduate student researcher,	Jun–Aug 2018
	Jelinek Memorial Summer Workshop (JSALT)	
	<i>PIs: Sam Bowman, Ellie Pavlick</i>	
	General-purpose Sentence Representation Learning team	
	Visiting researcher,	2015–2016

AWARDS & FUNDING	Korea Advanced Institute of Science and Technology (KAIST)	
	<i>PI: Jong C. Park</i>	
	Computational modeling of semantic fluency	
	Intern developer, NAVER Corporation	2013–2014
	Knowledge Extraction Team	
	Research partnership, MassMutual (\$40,000)	2025
	<i>For Mechanistic understanding of reasoning and explanations in entity tracking</i>	
	Best Paper Award, GenBench	2024
	<i>For Is artificial intelligence still intelligence?</i>	
	<i>LLMs generalize to novel adjective-noun pairs, but don't mimic the full human population</i>	
	National Deep Inference Fabric (NDIF) Pilot Program	2024
	<i>For Mechanisms of entity tracking in language models (with Sebastian Schuster)</i>	
	Unrestricted gift, Google (\$60,000)	2024
	Benchmarking LLM agents on consequential real-world tasks	2024–2026
	Open Philanthropy (\$756,396) with Sebastian Schuster	
	PaliGemma Academic Program, Google (\$5,000 in GCP credit)	2024
	<i>Comparing linguistic representations learned by text-only vs. multimodal models (with Kanishka Misra, Adam Dahlgren Lindström, Lucia Donatelli, & Yulu Qin)</i>	
	CDS Faculty Fellowship	2023–2026
	Center for Data Science, Boston University (\$50,000)	
	Area Chair Award (Interpretability and Analysis Track), ACL 2023	2023
	<i>For Entity Tracking in Language Models</i>	
	NSF Doctoral Dissertation Research Improvement Grant:	2021–2022
	<i>Compositional Linguistic Generalization in Human and Machine Learning</i>	
	BCS-2041221 (\$12,771) with Paul Smolensky, Geraldine Legendre & Tal Linzen.	
	NeurIPS 2019 Travel Award	2019
	Best Paper Award, *SEM 2019	2019
	<i>For Probing what different NLP tasks teach machines about function word comprehension</i>	
	Owen Scholarship	2016–2019
	Awarded to select incoming doctoral students in natural sciences (\$18,000)	
	Mica and Ahmet Ertegun Graduate Scholarship Programme in the Humanities	2014–2015
	Full scholarship for Master's degree at University of Oxford (~\$42,000)	
	SNU Undergraduate Research Program	2013
	Grant for select undergraduate research proposals (~\$3000)	
	SNU Foundation Funds Scholarship	2013
	SNU Superior Academic Performance Scholarship	2010–2011

WORK IN PROGRESS

Nicholas Edwards,* Yookyung Lee,* Yujun Audrey Mao, Yulu Qin, Sebastian Schuster,[†] **Najoung Kim**[†] (2025). [RExBench: Can coding agents autonomously implement AI research extensions?](#) In submission. (*,[†]Equal contribution)

Yulu Qin,* Dheeraj Varghese,* Adam Dahlgren Lindström, Lucia Donatelli, Kanishka Misra,[†] and **Najoung Kim**[†] (2025). [Vision-and-Language Training Helps Deploy Taxonomic Knowledge but Does Not Fundamentally Alter It](#). In submission. (*,[†]Equal contribution)

Yookyung Lee, JoongHoon Kim, Jaehee Kim, Hyowon Cho, Jaewook Kang, Pilsung Kang, and **Najoung Kim** (2025). [CheckEval: A reliable LLM-as-a-Judge framework for evaluating text generation using checklists](#). In submission.

Arkadiy Saakyan, Josh Lee, Michal Todorovic, Deepak Ramachandran, Quan Yuan, Isabelle Guyon, and **Najoung Kim** (2024). Evaluating Critic Models for Human-AI Co-Creation: A Case Study with AI Critiques of Presentation Slides. In submission.

Kanishka Misra and **Najoung Kim** (2024). [Generating novel experimental hypotheses from language models: A case study on cross-dative generalization](#). Preprint.

PUBLICATIONS

Hayley Ross, Kathryn Davidson, and **Najoung Kim** (2025). [Is analogy enough to draw novel adjective-noun inferences?](#) *Society for Computation in Linguistics (SCiL)*. Selected as talk

Abulhair Saparov, Srushti Pawar, Shreyas Pimpalgaonkar, Nitish Joshi, Richard Yuanzhe Pang, Vishakh Padmakumar, Seyed Mehran Kazemi, **Najoung Kim**,* and He He* (2025). [Transformers Struggle to Learn to Search Without In-context Exploration](#). *International Conference on Learning Representations (ICLR)*. (*Equal contribution)

Hayley Ross, **Najoung Kim**, and Kathryn Davidson (2025). [Fake reefs are sometimes reefs and sometimes not, but are always compositional](#). *Experiments in Linguistic Meaning (ELM)* 3. Selected as talk

Aditya Yedetore and **Najoung Kim** (2024). [Semantic Training Signals Promote Hierarchical Syntactic Generalization in Transformers](#). *Empirical Methods in Natural Language Processing (EMNLP)*.

Nitish Joshi, Javier Rando, Abulhair Saparov, **Najoung Kim**, and He He (2024). [Persons as a Way to Model Truthfulness in Language Models](#). *Empirical Methods in Natural Language Processing (EMNLP)*.

Hayley Ross, Kathryn Davidson, and **Najoung Kim** (2024). [Is artificial intelligence still intelligence? LLMs generalize to novel adjective-noun pairs, but don't mimic the full human population](#). *GenBench Workshop @ EMNLP*. **Best Paper Award**

Najoung Kim,* Sebastian Schuster,* and Shubham Toshniwal* (2024). [Code Pre-training Improves Entity Tracking Abilities of Language Models](#). arXiv. (*Equal contribution)

Najoung Kim and Paul Smolensky (2024). [Structural Generalization of Modification in Adult Learners of an Artificial Language](#). *Proceedings of the Annual Meeting of the Cognitive Science Society (CogSci)*. Selected as talk

Katherine M. Collins, **Najoung Kim**, Yonatan Bitton, Verena Rieser, Shayegan Omid-

shafiei, Yushi Hu, Sherol Chen, Senjuti Dutta, Minsuk Chang, Kimin Lee, Youwei Liang, Georgina Evans, Sahil Singla, Gang Li, Adrian Weller, Junfeng He, Deepak Ramachandran, and Krishnamurthy Dj Dvijotham (2024). [Beyond Thumbs Up/Down: Untangling Challenges of Fine-Grained Feedback for Text-to-Image Generation](#). *The Seventh AAAI/ACM Conference on AI, Ethics, and Society (AIES)*.

Ashwin Daswani, Rohan Sawant, and **Najoung Kim** (2024). [Syn-\(QA\)²: Evaluating False Assumptions in Long-tail Questions with Synthetic QA Datasets](#). arXiv.

Zhaofeng Wu, Linlu Qiu, Alexis Ross, Ekin Akyürek, Boyuan Chen, Bailin Wang, **Najoung Kim**, Jacob Andreas, and Yoon Kim (2024). [Reasoning or Reciting? Exploring the Capabilities and Limitations of Language Models Through Counterfactual Tasks](#). *North American Chapter of the Association for Computational Linguistics (NAACL)*.

Bingzhi Li, Lucia Donatelli, Alexander Koller, Tal Linzen, Yuekun Yao, and **Najoung Kim** (2023). [SLOG: A Structural Generalization Benchmark for Semantic Parsing](#). *Empirical Methods in Natural Language Processing (EMNLP)*.

Jason Wei,* **Najoung Kim**,* Yi Tay, and Quoc V. Le (2023). [Inverse scaling can become U-shaped](#). *Empirical Methods in Natural Language Processing (EMNLP)*. (*Equal contribution)

Najoung Kim* and Sebastian Schuster* (2023). [Entity Tracking in Language Models](#). *Annual Conference of the Association for Computational Linguistics (ACL)*. (*Equal contribution) **Area Chair Award**

Najoung Kim*, Phu Mon Htut*, Samuel R. Bowman, and Jackson Petty (2023). [\(QA\)²: Question Answering with Questionable Assumptions](#). *Annual Conference of the Association for Computational Linguistics (ACL)*. (*Equal contribution)

Seyed Mehran Kazemi, **Najoung Kim**, Deepti Bhatia, Xin Xu, and Deepak Ramachandran (2023). [LAMBADA: Backward Chaining for Automated Reasoning in Natural Language](#). *Annual Conference of the Association for Computational Linguistics (ACL)*.

Najoung Kim, Jatin Khilnani, Alex Warstadt, and Abed Qaddoumi (2023). [Reconstruction Probing](#). *Findings of the Annual Conference of the Association for Computational Linguistics (Findings of ACL)*.

Abulhair Saparov, Richard Yuanzhe Pang, Vishakh Padmakumar, Nitish Joshi, Seyed Mehran Kazemi, **Najoung Kim**,* and He He* (2023). [Testing the General Deductive Reasoning Capacity of Large Language Models Using OOD Examples](#). *Conference on Neural Information Processing Systems (NeurIPS)*. (*Equal contribution)

Mehran Kazemi, Quan Yuan, Deepti Bhatia, **Najoung Kim**, Xin Xu, Vaiva Imbrasaite, and Deepak Ramachandran (2023). [BoardgameQA: A Dataset for Natural Language Reasoning with Contradictory Information](#). *Conference on Neural Information Processing Systems (NeurIPS), Datasets and Benchmark Track*.

Ian McKenzie, Alexander Lyzhov, Michael Pieler, Alicia Parrish, Ameya Prabhu, Aaron Mueller, Euan McLean, Aaron Kirtland, Alexis Ross, Alisa Liu, Andrew Gritsevskiy, Daniel Wurgaft, Derik Kauffman, Gabriel Recchia, Jiacheng Liu, Joe Cavanagh, Max Weiss, Sicong Huang, The Floating Droid, Tom Tseng, Tomasz Korbak, Xudong Shen, Yuhui Zhang, Zhengping Zhou*, **Najoung Kim**, Samuel R. Bowman,

and Ethan Perez (2023). Inverse Scaling: When Bigger Isn't Better. *Transactions on Machine Learning Research (TMLR)*. (*Winning task authors) **Featured Certification**

Wentao Wang, Wai Keen Vong, **Najoung Kim**, and Brenden Lake (2023). [Finding Structure in One Child's Linguistic Experience](#). *Cognitive Science*.

Najoung Kim, Tal Linzen, and Paul Smolensky (2022). [Uncontrolled Lexical Exposure Leads to Overestimation of Compositional Generalization in Pretrained Models](#). arXiv.

Najoung Kim and Paul Smolensky (2021). [Testing for Grammatical Category Abstraction in Neural Language Models](#). *The Society for Computation in Linguistics (SCiL)*. [Selected as talk](#)

Najoung Kim, Ellie Pavlick, Burcu Karagol Ayan, and Deepak Ramachandran (2021). [Which Linguist Invented the Lightbulb? Presupposition Verification for Question-Answering](#). *Annual Conference of the Association for Computational Linguistics (ACL)*.

Najoung Kim and Tal Linzen (2020). [COGS: A Compositional Generalization Challenge Based on Semantic Interpretation](#). *Empirical Methods in Natural Language Processing (EMNLP)*.

Najoung Kim, Song Feng, Chulaka Gunasekara, and Luis A. Lastras (2020). [Implicit Discourse Relation Classification: We Need to Talk About Evaluation](#). *Annual Conference of the Association for Computational Linguistics (ACL)*.

Najoung Kim and Tal Linzen (2019). [Compositionality as directional consistency in sequential neural networks](#). *Context and Compositionality in Biological and Artificial Neural Systems Workshop, NeurIPS*.

Najoung Kim, Roma Patel, Adam Poliak, Alex Wang, Patrick Xia, Tom McCoy, Ian Tenney, Alexis Ross, Tal Linzen, and Benjamin Van Durme, Samuel R. Bowman, and Ellie Pavlick (2019). [Probing what different NLP tasks teach machines about function word comprehension](#). *Proceedings of the Eighth Joint Conference on Lexical and Computational Semantics (*SEM)*. **Pink Paper Award**

Alex Wang, Jan Hula, Patrick Xia, Raghavendra Pappagari, R. Thomas McCoy, Roma Patel, **Najoung Kim**, Ian Tenney, Yinghui Huang, Katherin Yu, Shuning Jin, Berlin Chen, Benjamin Van Durme, Edouard Grave, Ellie Pavlick, and Samuel R. Bowman (2019). [How to Get Past Sesame Street: Sentence-Level Pretraining Beyond Language Modeling](#). *Proceedings of the Annual Conference of the Association for Computational Linguistics (ACL)*.

Najoung Kim, Jung-Ho Kim, Maria K. Wolters, Sarah E. MacPherson, and Jong C. Park (2019). [Automatic Scoring of Semantic Fluency](#). *Frontiers in Psychology*.

Najoung Kim, Kyle Rawlins, Benjamin Van Durme, and Paul Smolensky (2019). [Predicting the Argumenthood of English Prepositional Phrases](#). *Proceedings of the 33rd AAAI Conference on Artificial Intelligence (AAAI)*.

Ian Tenney, Patrick Xia, Berlin Chen, Alex Wang, Adam Poliak, R. Thomas McCoy, **Najoung Kim**, Benjamin Van Durme, Samuel R. Bowman, Dipanjan Das, and Ellie Pavlick (2019). [What do you learn from context? Probing for sentence structure in contextualized word representations](#). *International Conference on Learning Representa-*

tions (ICLR).

Maria K. Wolters, **Najoung Kim**, Jung-Ho Kim, Sarah E. MacPherson, and Jong C. Park (2016). [Prosodic and Linguistic Analysis of Semantic Fluency Data: A Window into Speech Production and Cognition](#). *Interspeech*.

Jung-Ho Kim, **Najoung Kim**, Hancheol Park, and Jong C. Park (2016). Enhanced Sign Language Transcription System via Hand Tracking and Pose Estimation. *Journal of Computing Science and Engineering* vol 10.3.

Najoung Kim and Jong C. Park (2016). [A Morphological Approach to the Longitudinal Detection of Dementia](#). *Proceedings of HCI Korea 2016, The HCI Society of Korea*.

PRESENTATIONS WITHOUT PROCEEDINGS

Aditya Yedetore and **Najoung Kim** (2025). Implicit mechanisms for symbol manipulation in RNNs. *New England NLP Meeting Series (NENLP)*.

Hayley Ross, Kathryn Davidson, and **Najoung Kim** (2025). Is analogy enough to draw novel adjective-noun inferences? *New England NLP Meeting Series (NENLP)*.

Aditya Yedetore and **Najoung Kim** (2024). Semantic Training Signals Promote Hierarchical Syntactic Generalization in Neural Networks.

- *The 49th Boston University Conference on Language Development (BUCLD)*.
- *New England NLP Meeting Series (NENLP)*. [Selected as spotlight talk](#)

Hayley Ross, Kathryn Davidson, and **Najoung Kim** (2024). When is a fake concert still a concert? A study of adjective-noun composition in LLMs. *New England NLP Meeting Series (NENLP)*.

Kanishka Misra and **Najoung Kim** (2023). [Abstraction via exemplars? A representational case study on lexical category inference in BERT](#). Presented at *The 48th Boston University Conference on Language Development (BUCLD)*.

Sadhwi Srinivas,* **Najoung Kim**,* and Kyle Rawlins (2020). [Maximize Presupposition! and the Korean demonstrative ku](#). Presented at *The 94th Annual Meeting of the Linguistic Society of America (LSA)*. [[poster](#)] (*Equal contribution)

Najoung Kim, Benjamin Van Durme, Ellie Pavlick, and Paul Smolensky (2018). Linguistically informed tasks for evaluating structure encoded by sentence representations. *WeCNLP Summit*, Facebook HQ, Menlo Park, CA. [Selected as spotlight talk](#)

Najoung Kim, Kyle Rawlins, and Paul Smolensky (2018). A gradient blend analysis of English PP verbal dependents. *Conference on Interdisciplinary Approaches to Linguistic Theory (CiALT) 2*, Humboldt-Universität zu Berlin.

Najoung Kim, Kyle Rawlins, and Paul Smolensky (2018). A gradient blend analysis of English PP verbal dependents. *Acceptability judgments in current linguistic theory*, Universitat Autònoma de Barcelona.

OTHER ARTICLES

Yonatan Belinkov, **Najoung Kim**, Jaap Jumelet, Hosein Mohebbi, Aaron Mueller, and Hanjie Chen (2024). [Proceedings of the 7th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP](#). Association for Computational Linguistics.

Yonatan Belinkov, Sophie Hao, Jaap Jumelet, **Najoung Kim**, Arya McCarthy, and Hosein Mohebbi (2023). [Proceedings of the 6th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP](#). Association for Computational Linguistics.

Teven Le Scao, Angela Fan, ... **Najoung Kim** (Evaluation and Interpretability) ... (2022), [BLOOM: A 176B-Parameter Open-Access Multilingual Language Model](#). arXiv.

Najoung Kim (2022). [Compositional Generalization in Artificial Neural Networks](#). PhD Dissertation, Johns Hopkins University. Committee: Paul Smolensky, Kyle Rawlins, Benjamin Van Durme, Tal Linzen, and Bob Frank.

SELECTED TALKS / TUTORIALS *What does it take to convince ourselves that a system is exhibiting compositionality?*

- Keynote at [Copenhagen NLP Symposium](#). June 2025.
- Keynote at the [10th Workshop on Representation Learning for NLP \(RepL4NLP\)](#), co-located with NAACL 2025. May 2025.

Evaluating Multimodal Critic Models for Human-AI Co-Creation. Keynote at the [Structure and Generalization in Multimodal Language Understanding \(SAGE-MLU\)](#) Workshop. March 2025.

Semantic generalizations in humans and machines.

- Keynote at the [Workshop on Generalization in NLP \(GenBench\) 2024](#) (co-located with EMNLP 2024). November 2024.
- Invited talk at Sungkyunkwan University, December 2024.

Generating novel experimental hypotheses from language models: A case study on cross-dative generalization. Invited talk at the [Language & Cognition Seminar at Harvard \(LangCog\)](#). October 2024.

Human and Machine Inductive Biases for Compositional Linguistic Generalization. Invited talk at the [Emerging Generalization Settings Workshop at Simons Institute](#). September 2024.

Comparing human and machine inductive biases for compositional linguistic generalization using semantic parsing. Invited talk at the [Fifth International Convention on the Mathematics of Neuroscience and AI \(Neuromonster\)](#). May 2024.

Linguistic tests as unit tests for AI systems. Invited talk at [New Horizons in Language Science: Large Language Models, Language Structure, and the Cognitive and Neural Basis of Language](#). National Science Foundation (NSF). May 2024.

Entity Tracking in Language Models.

- Invited talk at [UMass NLP](#). November 2024.
- Invited talk at KAIST Electrical Engineering. July 2024.
- Invited talk at AIR seminar at Boston University. April 2024.
- Invited talk at [South by Semantics Workshop at UT Austin](#). Feb 2024.
- Invited talk at [Penn NLP CLunch Seminar](#). Oct 2023.

Can we use simulated learners to prune the hypothesis space for human experiments?

- Invited talk at CPL Lab, MIT. April 2024.

- Invited talk at [Brown Linguistics](#). March 2024.

Inverse scaling: mitigation strategies and open questions. Keynote talk at the [Workshop on the Scaling Behavior of Large Language Models \(SCALE-LLM\)](#). March 2024.

Challenges in logical reasoning with LLMs. Invited talk at [Yale/Google ML Workshop: Theory and Practice of Foundation Models](#). Oct 2023.

Evidence for abstraction-via-exemplars from lexical category inference in neural language models. Invited talk at the [Virtual Psycholinguistics Forum](#). Sep 2023.

Compositional Linguistic Generalization in Artificial Neural Networks: Taking Stock. Invited talk at [Harvard Universals Workshop](#), Apr 2023.

Questions with Questionable Assumptions as a Challenge to QA Systems. Invited talk at Colgate University. Mar 2023.

Compositional Linguistic Generalization in Artificial Neural Networks: Taking Stock.

- Invited talk at Seminars on Formal Languages and Neural Networks (FlaNN), Nov 2022.
- Invited talk at CUNY Graduate Center Computational Linguistics Talk Series, Nov 2022.

Compositional Linguistic Generalization in Artificial Neural Networks.

- Invited talk at CompLing Lab, University of Chicago, May 2022.
- ‘Tech Talk’ at Google, May 2022.
- Invited talk at Boston University Linguistics Colloquium Series, Mar 2022.
- Invited talk at [MIT Complang](#). Mar 2022.
- Invited talk at [Seminar on the interactions between formal and computational linguistics \(ILFC\)](#). Feb 2022.
- Invited talk at Linguistics Colloquium, Seoul National University. Dec 2021.
- Invited talk at [KAIST Electrical Engineering Colloquium Lecture Series](#). Nov 2021.
- Invited talk at [NYU ConCats](#). Nov 2021.
- Invited talk at the University of Michigan Cognitive Science Community Colloquium. Oct 2021.
- Invited talk at the [New York Philosophy of Language Workshop](#). Oct 2021.
- Invited talk at NYU Center for Data Science. Sep 2021.

Compositional Linguistic Generalization in Contemporary Neural Models of Language. Invited talk at Cornell Computational Psycholinguistics Group. Apr 2021.

What Aspects of Meaning are Missing from Current Natural Language Understanding Systems? Invited talk at Boston University, Mar 2021.

COGS: A Compositional Generalization Challenge Based on Semantic Interpretation.

- Invited talk at NERT lab, Georgetown University, Nov 2020.
- Invited poster presentation at the [Microsoft Research AI Breakthroughs Workshop](#), Sep 2020.

Probing what different NLP tasks teach machines about function word comprehension and where to go next. Invited talk at JHU CLSP Seminar, Oct 2019.

The complement-adjunct distinction as gradient blends: the case of English prepositional phrases. Invited talk at the Gradient Symbolic Computation Workshop, Baltimore, Sep 2019.

Semantic Role Labeling Tutorial (with Diego Marcheggiani, Michael Roth, and Benjamin Van Durme), EMNLP, Copenhagen, Sep 2017.

Approximating the Semantic Structures behind Category Fluency Sequences (Poster). MACSIM 6 at CUNY, New York, Oct 2016.

Detection and Categorisation of Neograms in Korean Text (+a). Invited talk at Ertegun House, Oxford, 2015.

MEDIA COVERAGE

[Linguistics in the Age of Artificial Intelligence](#). Boston University Arts & Sciences Magazine, April 2024.

[Compositional Generalization in Neural Networks](#). Interview at NLP Highlights Podcast hosted by Allen Institute for AI, Jan 2023.

ADVISING / MENTORING

Postdoctoral Fellows:

- [Yukyung Lee](#), 2024–

PhD Advisees:

- [Yulu Qin](#) (BU Linguistics), 2024–
- [Audrey Mao](#) (BU Computer Science), 2024–
- [Hayley Ross](#) (Harvard Linguistics, co-advised with Kathryn Davidson), 2023–2025 → *NVIDIA*
- [Jing Liu](#) (École Normale Supérieure, co-advised with Emmanuel Dupoux & David Harwarth), 2023–
- [Aditya Yedetore](#) (BU Linguistics), 2022–

PhD Committees:

- Ousmane Cisse (Dissertation prospectus, BU Linguistics), 2025–
- Cheng Tang (Dissertation, MIT Brain & Cognitive Science), 2025–
- Afra Feyza Akyürek (Dissertation & Prospectus, BU Computer Science), 2024–2025
- Yusuf Kocyigit (Qualifying exam, BU Computer Science), 2025
- Gabriel Franco (Qualifying exam, BU Computer Science), 2025
- Zhongping Zhang (Dissertation & Prospectus, BU Computer Science), 2024
- Andrea Burns (Dissertation, BU Computer Science), 2023
- Isidora Tourni (Dissertation, BU Computer Science), 2023

Other Ad-hoc Advising & Mentoring:

- Xiang Fu, Supported by [BU Undergraduate Research Opportunities Program \(UROP\)](#), Summer 2025

- Nicholas Edwards (co-mentored with Sebastian Schuster), Research Assistant at University College London. 2024–
- Arkadiy Saakyan (co-hosted with Deepak Ramachandran), Student Researcher at Google DeepMind, May–Aug 2024.
 - Outcome: Paper submitted to DIS
- Alyssa Feinberg (undergraduate at Brown University), Research Assistant @ BU, Summer 2024.
- Alara Balcisoy, Undergraduate researcher @ BU, Spring 2024.
 - Supported by [BU Undergraduate Research Opportunities Program \(UROP\)](#)
 - Outcome: Recipient of [BUCH Award](#)
- Bingzhi Li (co-advised with Tal Linzen), Visiting PhD student from Université de Paris, Fall 2022.
 - Outcome: Paper at EMNLP
- Abed Qaddoumi (co-mentored with Alex Warstadt), MS at New York University. 2021–2022. *Current: PhD student at Stony Brook Linguistics*
 - Outcome: Paper at ACL Findings
- Jatin Khilnani (co-mentored with Alex Warstadt), MS at New York University. 2021–2022. *Current: PhD student at University of Pittsburgh CS*
 - Outcome: Paper at ACL Findings
- Pablo Santos (co-mentored with Phu Mon Htut), [CDS-Courant Undergraduate Research Program \(CURP\)](#), Spring 2022.

TEACHING

Boston University:

- Fall 2025, Topics in Linguistics: Contemporary Research in Neural Network Models of Language
- Fall 2024–2025, Topics in Linguistics: Cognitive Science of Language
- Spring 2024–2025, Computational Linguistics
- Fall 2023–2024, Introduction to Programming for Computational Linguistics
- Spring 2023–2024, Topics in Linguistics: Metrics and Evaluation in Natural Language Processing

New York University:

- Fall 2021–2022, Capstone Project in Data Science

Johns Hopkins University:

- Fall 2019, Introduction to Computational Cognitive Science (TA/Co-Instructor)
- Spring 2019, Foundations of Neural Networks (TA/Lab Instructor)
- Spring 2018, Foundations of Cognitive Science (TA)
- Fall 2017, Semantics I (TA/Lab Instructor)
- Spring 2017, Language and Advertising (TA)

Seoul National University:

- Winter 2014, Samsung Convergence Software Course (Instructor)

PROFESSIONAL SERVICE

Grant panels:

- Advisory Panel, US National Science Foundation (NSF) (2024)

Conference Chairs:

- ACL 2026 (Website Chair)
- EMNLP 2025 (Area Chair)
- COLM 2025 (Area Chair)
- ACL 2025 (Senior Area Chair for Linguistic Theories, Cognitive Modeling, and Psycholinguistics)
- EMNLP 2024 (Senior Area Chair for Linguistic Theories, Cognitive Modeling, and Psycholinguistics)
- NAACL 2024 (Publicity Chair)
- ACL 2023 (Senior Area Chair for Linguistic Theories, Cognitive Modeling, and Psycholinguistics)
- *SEM 2022 (Area Chair for Psycholinguistics, cognitive linguistics and semantic processing)

Organization:

- [New England Mechanistic Interpretability Workshop](#) (2025)
- [BlackboxNLP 2025](#) (Co-located with EMNLP 2025)
- [BlackboxNLP 2024](#) (Co-located with EMNLP 2024)
- AI4Research (Co-located with IJCAI 2024)
- [BlackboxNLP 2023](#) (Co-located with EMNLP 2023)
- [Inverse Scaling Prize](#) (2022)

Ad-hoc journal reviews:

- Journal of Memory and Language (2024)
- Computational Linguistics (2024)
- Language, Cognition and Neuroscience (2024)
- Cognitive Science (2023, 2022)
- Nature Machine Intelligence (2023, 2022)

Ad-hoc grant reviews:

- Israel Science Foundation (2024)

Conference/Workshop reviews:

- NAACL Student Research Workshop (participated as mentor, 2025)
- COLM (2024)
- *SEM (2024, 2023, * 2022, 2019)
- Scale-LLM Workshop (2024)
- EACL Student Research Workshop (2023, 2021)
- EMNLP (2022, 2020, * 2018)
- COLING (2022)

- Bridges and Gaps between Formal and Computational Linguistics (Workshop at ESSLLI 2022)
- CogSci (2022)
- ACL Student Research Workshop, (2019–2022)
- ACL-IJCNLP (2021)
- ACL (2020)
- EMNLP-IJCNLP (2019)

*Acknowledged as outstanding reviewer

PROFESSIONAL MEMBERSHIP	Association for Computational Linguistics Cognitive Science Society Linguistic Society of America
COMPUTER SKILLS	Python, R, HTML (proficient) Java, C, Scheme/Racket, OCaml (basic) Tools: PyTorch, PCIBex, Mechanical Turk, Prolific
HUMAN LANGUAGES	Korean (native); English (near-native); Spanish (intermediate); Japanese (reading knowledge)
COMMUNITY ENGAGEMENT	<p>Dagstuhl Seminar, Linguistics and Language Models: What Can They Learn from Each Other? 2025</p> <p>Panelist, New England NLP Meeting Series (NENLP) 2025</p> <p>Panelist, Learning to Think after ChatGPT: A Panel Discussion <i>Center for Data Science, Boston University</i> 2023</p> <p>Panelist, KASELL Fall Conference 2021</p> <p>Diversity & Representation Committee 2021 Colloquium Student Committee 2018–2021 <i>Department of Cognitive Science, Johns Hopkins University</i></p> <p>Panelist, BrainIAC Professional Development Event 9/21/2020</p> <p>Organizer, Corpus Tools: Past, Present, and Future. 2015 <i>Ertegun House, Oxford</i></p> <p>Mentor, Samsung Convergence Software Course 2014 <i>Department of Computer Science, Seoul National University</i></p> <p>Peer tutor, English (volunteer work) 2011 <i>Seoul National University</i></p> <p>English tutor (volunteer work) 2010 <i>Yongnak High School</i></p>